

Ondřej Bílek
ondrej@h2o.ai

Introduction to Enterprise Steam with Sparkling Water

H2O.ai, Prague
Oct 16, 2018

About me

- System programming @ CTU in Prague
- Software Engineer @ H2O.ai
- Codebase owner of Enterprise Steam

The Hadoop problem

- Hadoop configuration is hard!
- Existing tools focus on admins, not users
- Multi-user environments are tricky
- Limited control over resources
- Security over convenience

H₂O Products



H₂O.ai

In-Memory, Distributed
Machine Learning Algorithms
with H2O Flow GUI



Spark + H₂O
SPARKLING
WATER

H2O AI Open Source Engine
Integration with Spark



H₂O4GPU

Lightning Fast machine
learning on GPUs



DRIVERLESSAI

Automatic feature
engineering, machine learning
and interpretability


Enterprise Steam

Secure multi-tenant H2O clusters

Enterprise Steam

- Multi-tenant system for managing H2O clusters on Hadoop
- Powerful admin capabilities
- Simplicity in mind
- Web frontend, Python & R API
- Enterprise grade security

Steam + Sparkling Water



In-Memory, Distributed
Machine Learning Algorithms
with H2O Flow GUI



H2O AI Open Source Engine
Integration with Spark



Lightning Fast machine
learning on GPUs



Automatic feature
engineering, machine learning
and interpretability

Enterprise Steam

Secure multi-tenant H2O clusters

Why?

- Setting-up SW on Hadoop is not trivial
- Restricted to Hadoop node
- Reusable Spark sessions
- Pre-made configurations / templates
- Notebook integration (Jupyter)
- Security

Demo Time

How it works?

